

DISTRIBUTED AI SYSTEMS

Build scalable training, inference and serving systems

With large AI models now reaching billions or even trillions of parameters, distributed systems have become essential for training and serving AI. This comprehensive guide bridges the gap between theory and practice with hands-on code examples demonstrating production-grade techniques used in real-world systems, from distributed training through inference to production serving.

Key Features

- Master distributed training with DDP, FSDP, DeepSpeed, and Megatron
- Build high-performance inference systems with vLLM and SGLang
- Understand GPU hardware, interconnects, and parallelism strategies
- Deploy production serving stacks with orchestration and observability
- Hands-on code examples tested on real infrastructure throughout

Intended Audience

Designed for ML engineers, AI researchers, and DevOps engineers who need to train or serve large AI models at scale. Platform engineers, HP Cluster administrators, and cloud architects will advance their skills with this guide.